

naytzwqed

December 4, 2024

Question 1 - LSTM Autoencoder You are required to build an LSTM Autoencoder to detect anomalies in a time series dataset. The dataset contains daily temperature readings from a weather station over the course of a few years. Your task is to: 1. Load the dataset: The dataset will contain a single column temperature and a date column. 2. Preprocess the data: Normalize the temperature data and split it into training and testing sets. 3. Build an LSTM Autoencoder: o The encoder should reduce the input dimensions to a latent representation. o The decoder should reconstruct the input from the latent representation. 4. Train the model: Train the autoencoder on the training data and evaluate the reconstruction error on the test set. 5. Anomaly Detection: Use the reconstruction error to detect anomalies. Define a threshold for the reconstruction error, and identify days where the temperature is considered anomalous. 6. Visualize the results: Plot the original temperature data and highlight the detected anomalies.

1. Load the weather\_data.csv - Dataset and Parse the Date column as datetime and ensure the Temperature column is numeric.

```
[54]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from tensorflow.keras.layers import Input, LSTM, RepeatVector, Dense
from tensorflow.keras.models import Model
```

```
[55]: def load_data(filepath='/content/weather_data.csv'):
    df = pd.read_csv(filepath)
    df['date'] = pd.to_datetime(df['date'])
    return df.set_index('date')
```

2. Exploratory Data Analysis

```
[56]: df.tail()
```

```
[56]:
```

	date	temperature
2552	2020-12-27	10.598482
2553	2020-12-28	9.876287
2554	2020-12-29	9.265936
2555	2020-12-30	9.986115
2556	2020-12-31	8.419751

```
[57]: df.describe()
```

```
[57]:      temperature
count  2557.000000
mean    10.017472
std      0.923047
min      3.337291
25%      9.335195
50%     10.031778
75%     10.681384
max     15.587945
```

```
[58]: df.head()
```

```
[58]:      date  temperature
0  2014-01-01    10.248357
1  2014-01-02     9.950428
2  2014-01-03    10.362958
3  2014-01-04    10.820167
4  2014-01-05     9.961091
```

```
[59]: df.isnull().sum()
```

```
[59]: date          0
      temperature  0
      dtype: int64
```

```
[73]: def perform_eda(df):

    plt.figure(figsize=(15, 12))

    # Time Series Plot
    plt.subplot(3, 1, 1)
    plt.plot(df.index, df['temperature'])
    plt.title('Temperature Time Series')
    plt.xlabel('Date')
    plt.ylabel('Temperature')

    # Distribution Plot
    plt.subplot(3, 1, 2)
    sns.histplot(df['temperature'], bins=30, kde=True)
    plt.title('Temperature Distribution')

    # Monthly Box Plot
    plt.subplot(3, 1, 3)
    df['month'] = df.index.month
    sns.boxplot(x='month', y='temperature', data=df)
```

```
plt.title('Monthly Temperature Distribution')

plt.tight_layout()
plt.show()

# Print basic statistics
print("\nBasic Statistics:")
print(df['temperature'].describe())
```

### 3. Model Implementation

```
[61]: def create_sequences(data, seq_length=7):
        return np.array([data[i:i+seq_length] for i in
        ↪range(len(data)-seq_length+1)])
```

Create the AutoEncoder

```
[70]: def build_autoencoder(seq_length=7):
        # Encoder
        inputs = Input(shape=(seq_length, 1))
        encoded = LSTM(32, activation='relu')(inputs)

        # Decoder
        decoded = RepeatVector(seq_length)(encoded)
        decoded = LSTM(32, activation='relu', return_sequences=True)(decoded)
        decoded = Dense(1)(decoded)

        model = Model(inputs, decoded)
        model.compile(optimizer='adam', loss='mse')
        return model
```

Train The AutoEncoder Model

```
[72]: def train_model(model, X_train, X_test):
        history = model.fit(
            X_train, X_train,
            epochs=50,
            batch_size=32,
            validation_data=(X_test, X_test),
            verbose=1
        )
        return history
```

### 4. Results Analysis and Visualization

```
[81]: def analyze_results(model, df, test_data, test_sequences, scaler):
        """Analyze and visualize results"""
        # Get predictions and errors
```

```

predictions = model.predict(test_sequences)
mse = np.mean(np.square(test_sequences - predictions), axis=(1,2))
threshold = np.mean(mse) + 2 * np.std(mse)
anomalies = mse > threshold

# Plot results
plt.figure(figsize=(15, 10))

plt.subplot(2, 1, 1)
plt.plot(test_data.index, scaler.inverse_transform(test_data),
↪label='Original')
anomaly_points = test_data[anomalies]
plt.scatter(anomaly_points.index,
            scaler.inverse_transform(anomaly_points),
            color='red', label='Anomaly')
plt.title('Temperature Data with Detected Anomalies')
plt.legend()

plt.subplot(2, 1, 2)
plt.plot(mse, label='Reconstruction Error')
plt.axhline(y=threshold, color='r', linestyle='--', label='Threshold')
plt.title('Reconstruction Error and Threshold')
plt.legend()

plt.tight_layout()
plt.show()

#Insights
print("\nAnalysis Insights:")
print(f"Total anomalies detected: {sum(anomalies)}")
print(f"Anomaly percentage: {(sum(anomalies)/len(anomalies))*100:.2f}%")

return anomalies, mse, threshold

```

## 5. Main Funtion Executions

```

[84]: import pandas as pd
from sklearn.preprocessing import MinMaxScaler

def main():
    df = load_data()

    # 1. Perform EDA
    print("Performing Exploratory Data Analysis...")

```

```

perform_eda(df)

# 2. Preprocess data
scaler = MinMaxScaler()
scaled_data = scaler.fit_transform(df[['temperature']])

# Create sequences
seq_length = 7
sequences = create_sequences(scaled_data)

# Split data
train_size = int(len(sequences) * 0.8)
X_train = sequences[:train_size]
X_test = sequences[train_size:]
test_data = pd.DataFrame(
    scaled_data.flatten()[train_size + seq_length - 1:],
    index=df.index[train_size + seq_length - 1:],
    columns=['temperature']
)

# 3. Build and train model
print("\nTraining LSTM Autoencoder...")
model = build_autoencoder(seq_length)
history = train_model(model, X_train, X_test)


# 4. Analyze results
print("\nAnalyzing Results...")
anomalies, mse, threshold = analyze_results(model, df, test_data, X_test,
↪scaler)

return model, history, anomalies

if __name__ == "__main__":
    model, history, anomalies = main()

```

Performing Exploratory Data Analysis...



#### Basic Statistics:

```
count    2557.000000
mean      10.017472
std       0.923047
min       3.337291
25%      9.335195
50%      10.031778
75%      10.681384
max      15.587945
Name: temperature, dtype: float64
```

#### Training LSTM Autoencoder...

Epoch 1/50

64/64 4s 20ms/step -

loss: 0.1744 - val\_loss: 0.0212

Epoch 2/50

64/64 1s 20ms/step -

loss: 0.0149 - val\_loss: 0.0049

Epoch 3/50  
64/64 2s 9ms/step - loss:  
0.0038 - val\_loss: 0.0018  
Epoch 4/50  
64/64 1s 8ms/step - loss:  
0.0021 - val\_loss: 0.0017  
Epoch 5/50  
64/64 1s 9ms/step - loss:  
0.0020 - val\_loss: 0.0017  
Epoch 6/50  
64/64 1s 9ms/step - loss:  
0.0020 - val\_loss: 0.0017  
Epoch 7/50  
64/64 1s 9ms/step - loss:  
0.0020 - val\_loss: 0.0017  
Epoch 8/50  
64/64 1s 9ms/step - loss:  
0.0020 - val\_loss: 0.0017  
Epoch 9/50  
64/64 1s 9ms/step - loss:  
0.0020 - val\_loss: 0.0017  
Epoch 10/50  
64/64 1s 9ms/step - loss:  
0.0019 - val\_loss: 0.0017  
Epoch 11/50  
64/64 1s 9ms/step - loss:  
0.0019 - val\_loss: 0.0016  
Epoch 12/50  
64/64 1s 8ms/step - loss:  
0.0019 - val\_loss: 0.0016  
Epoch 13/50  
64/64 1s 9ms/step - loss:  
0.0019 - val\_loss: 0.0016  
Epoch 14/50  
64/64 1s 9ms/step - loss:  
0.0019 - val\_loss: 0.0016  
Epoch 15/50  
64/64 1s 8ms/step - loss:  
0.0019 - val\_loss: 0.0016  
Epoch 16/50  
64/64 1s 9ms/step - loss:  
0.0019 - val\_loss: 0.0016  
Epoch 17/50  
64/64 1s 8ms/step - loss:  
0.0019 - val\_loss: 0.0016  
Epoch 18/50  
64/64 1s 9ms/step - loss:  
0.0019 - val\_loss: 0.0016

Epoch 19/50  
64/64 1s 12ms/step -  
loss: 0.0019 - val\_loss: 0.0016  
Epoch 20/50  
64/64 1s 14ms/step -  
loss: 0.0019 - val\_loss: 0.0016  
Epoch 21/50  
64/64 1s 14ms/step -  
loss: 0.0018 - val\_loss: 0.0016  
Epoch 22/50  
64/64 1s 14ms/step -  
loss: 0.0018 - val\_loss: 0.0016  
Epoch 23/50  
64/64 1s 9ms/step - loss:  
0.0018 - val\_loss: 0.0016  
Epoch 24/50  
64/64 1s 9ms/step - loss:  
0.0018 - val\_loss: 0.0016  
Epoch 25/50  
64/64 1s 9ms/step - loss:  
0.0018 - val\_loss: 0.0016  
Epoch 26/50  
64/64 1s 9ms/step - loss:  
0.0018 - val\_loss: 0.0016  
Epoch 27/50  
64/64 1s 9ms/step - loss:  
0.0018 - val\_loss: 0.0016  
Epoch 28/50  
64/64 1s 9ms/step - loss:  
0.0018 - val\_loss: 0.0016  
Epoch 29/50  
64/64 1s 9ms/step - loss:  
0.0018 - val\_loss: 0.0016  
Epoch 30/50  
64/64 1s 9ms/step - loss:  
0.0018 - val\_loss: 0.0016  
Epoch 31/50  
64/64 1s 9ms/step - loss:  
0.0018 - val\_loss: 0.0016  
Epoch 32/50  
64/64 1s 9ms/step - loss:  
0.0018 - val\_loss: 0.0016  
Epoch 33/50  
64/64 1s 9ms/step - loss:  
0.0018 - val\_loss: 0.0015  
Epoch 34/50  
64/64 1s 10ms/step -  
loss: 0.0017 - val\_loss: 0.0015




Epoch 35/50  
64/64 1s 8ms/step - loss:  
0.0017 - val\_loss: 0.0015  
Epoch 36/50  
64/64 1s 9ms/step - loss:  
0.0017 - val\_loss: 0.0015  
Epoch 37/50  
64/64 1s 9ms/step - loss:  
0.0017 - val\_loss: 0.0015  
Epoch 38/50  
64/64 1s 11ms/step -  
loss: 0.0017 - val\_loss: 0.0015  
Epoch 39/50  
64/64 1s 14ms/step -  
loss: 0.0017 - val\_loss: 0.0015  
Epoch 40/50  
64/64 1s 13ms/step -  
loss: 0.0017 - val\_loss: 0.0014  
Epoch 41/50  
64/64 1s 8ms/step - loss:  
0.0017 - val\_loss: 0.0014  
Epoch 42/50  
64/64 1s 9ms/step - loss:  
0.0017 - val\_loss: 0.0014  
Epoch 43/50  
64/64 1s 9ms/step - loss:  
0.0016 - val\_loss: 0.0014  
Epoch 44/50  
64/64 1s 8ms/step - loss:  
0.0016 - val\_loss: 0.0014  
Epoch 45/50  
64/64 1s 9ms/step - loss:  
0.0016 - val\_loss: 0.0014  
Epoch 46/50  
64/64 1s 8ms/step - loss:  
0.0016 - val\_loss: 0.0013  
Epoch 47/50  
64/64 1s 9ms/step - loss:  
0.0016 - val\_loss: 0.0013  
Epoch 48/50  
64/64 1s 8ms/step - loss:  
0.0015 - val\_loss: 0.0013  
Epoch 49/50  
64/64 1s 9ms/step - loss:  
0.0015 - val\_loss: 0.0013  
Epoch 50/50  
64/64 1s 9ms/step - loss:  
0.0015 - val\_loss: 0.0013

Analyzing Results...

16/16

1s 23ms/step



Analysis Insights:

Total anomalies detected: 7

Anomaly percentage: 1.37%